

# Selective State Spaces

## Bridging the Gap Between RNNs and Transformers

Clement Marie • Ying JIN

Course: Deep Learning Project

## 1. Motivation & Goal

**The Problem:** While Transformers dominate deep learning, their quadratic complexity ( $O(N^2)$ ) limits scalability. Mamba (SSM) offers linear-time ( $O(N)$ ) inference, but official implementations are opaque (CUDA-heavy) and strictly causal (unidirectional).

### Our Goal:

1. Provide a transparent **Pure PyTorch** implementation of Selective Scan (S6).
2. Extend Mamba with **Bi-directional Inductive Bias** for Vision tasks.

## 2. Method: Mamba Architecture

Mamba replaces attention with an input-dependent SSM. Key components: **Learned Discretization  $\Delta$** , **Causal Conv1d**, and **Gating**.




Figure 1: (Top) Gated block with SiLU and residual; (Bottom) Selection mechanism where learned  $\Delta$  acts as a gate.

**From Continuous to Discrete:** Discretizing the continuous ODE  $h'(t) = \mathbf{A}h(t) + \mathbf{B}x(t)$  with a learned time-step  $\Delta$ :

$$h_t = \bar{\mathbf{A}}_t h_{t-1} + \bar{\mathbf{B}}_t x_t, \quad y_t = \mathbf{C}_t h_t$$

where  $\bar{\mathbf{A}}_t = \exp(\Delta_t \mathbf{A})$ ,  $\bar{\mathbf{B}}_t = (\Delta_t \mathbf{A})^{-1}(\bar{\mathbf{A}}_t - \mathbf{I})$ ,

*Innovation:*  $\Delta_t$  depends on input  $x_t$ , allowing the model to selectively "remember".

## 3. Algorithm: Selective Scan (S6)

The input-dependent recurrence is computed via a parallel scan.

### Algorithm 1 SSM + Selection (S6)

```

Require: Input  $x : (B, L, D)$ 
Ensure: Output  $y : (B, L, D)$ 
1:  $A : (D, N) \leftarrow$  Parameter Fixed structured matrix
   // 1. Selection: Projects input to parameters
2:  $B : (B, L, N) \leftarrow \text{Linear}_B(x)$ 
3:  $C : (B, L, N) \leftarrow \text{Linear}_C(x)$ 
4:  $\Delta : (B, L, D) \leftarrow \text{Softplus}(\text{Parameter} + \text{Linear}_\Delta(x))$ 
   // 2. Discretization
5:  $\bar{A}, \bar{B} \leftarrow \text{discretize}(\Delta, A, B)$ 
   // 3. Selective Scan (The parallel recurrence)
6:  $y \leftarrow \text{SSM}(\bar{A}, \bar{B}, C)(x)$ 
7: return  $y$ 

```

## 4. Implementation Details

- ▶ **Framework:** Pure PyTorch (No custom CUDA kernels).
- ▶ **Conv1d:** Kernel size 4, helps with local context.
- ▶ **Normalization:** RMSNorm for stability.
- ▶ **Optimization:** Weight tying between embeddings and output head.

The code can be found at :  
<https://github.com/clmrie/Mamba-S6-From-Scratch>.

## 5. Vision Extension: Bi-Directional

**Why Bi-directional?** Standard Mamba loses context from "future" pixels, which is suboptimal for images. Bi-directional scanning captures global spatial context. We implement a **Bi-Directional Vision Mamba**:

1. Flatten image into patch tokens.
2. Scan Forward ( $\rightarrow$ ) and Backward ( $\leftarrow$ ).
3. Fuse branches (Average/Concat).




Figure 2: Causal vs. Bi-Directional Performance.

## 6. Experimental Setup

- ▶ MNIST: Vision Mamba, causal ablation, vanilla RNN baseline.
- ▶ CIFAR-10: patch-based classification.
- ▶ TinyShakespeare: Character-level LM.

Dataset	Key training settings
MNIST (Vision)	d_model=64, layers=2, batch=64, lr=1e-3, epochs=5
MNIST (Causal)	d_model=64, layers=2, batch=64, lr=3e-4, epochs=5
MNIST (RNN)	input=16, hidden=64, layers=2, batch=64, lr=1e-3, epochs=5
CIFAR-10	d_model=128, layers=4, batch=64, lr=1e-3, epochs=5
TinyShakespeare	d_model=128, layers=4, block=128, batch=32, lr=1e-3, epochs=10

Table 1: Training settings summary

## 7. Results: Mamba vs RNN

Vision Mamba converges significantly faster than vanilla RNNs and achieves higher final accuracy.




Figure 3: Training Loss Comparison on MNIST.

Model Architecture	Acc. (%)
Vanilla RNN	94.40
Causal Vision Mamba	97.00
<b>Bi-Directional Vision Mamba (Ours)</b>	<b>97.93</b>

Table 2: Classification Accuracy.

## 7. Results: Visualization of delta selection

We analyze the magnitude of the learned time-step  $\Delta_t$  on a "Needle in a Haystack" task. The model exhibits distinct **activation spikes** (high  $\Delta$  values) precisely at the positions of critical information (the "needle"), while maintaining low values for irrelevant noise. Mathematically, a large  $\Delta_t$  increases the update rate of the state  $h_t$ , effectively "opening the gate" to write information into memory.




Figure 4: Visualization of  $\Delta$

## 8. Conclusion

- ▶ Successfully reimplemented **Mamba S6** from scratch without CUDA kernels.
- ▶ **Bi-directional scanning** is crucial for non-causal data (images), boosting accuracy by  $\sim 1\%$  over causal Mamba.
- ▶ Demonstrated linear-time scaling potential on sequence tasks.

## References

[1] Gu & Dao, *Mamba: Linear-Time Sequence Modeling with Selective State Spaces*, 2023.